

Machine Learning I (DATS 6202)


Linear Regression

Yuxiao Huang

Data Science, Columbian College of Arts & Sciences
George Washington University
yuxiaohuang@gwu.edu

September 10, 2018

Reference



Picture courtesy of the website of the book code repository and info resource

Reference

- This set of slides is an excerpt of the book by Raschka and Mirjalili, with some trivial changes by the creator of the slides
- Please find the reference to and website of the book below:
 - *Raschka S. and Mirjalili V. (2017). Python Machine Learning. 2nd Edition.*
 - <https://sebastianraschka.com/books.html>
- Please find the website of the book code repository and info resource below:
 - <https://github.com/rasbt/python-machine-learning-book-2nd-edition>

Overview

- 1 The linear regression model
- 2 Exploring and visualizing datasets
- 3 Evaluating regression models and diagnosing common problems
- 4 Implementing linear regression models
- 5 Training regression models that are robust to outliers

Simple (univariate) linear regression

- Simple linear regression expresses the relationship between two continuous-valued variables, x and y :


$$y = w_0 + w_1x$$

- Here:
 - x : explanatory variable (or independent variable, regressor)
 - y : response variable (or dependent variable, regressand)
 - w_0 : intercept
 - w_1 : slope
- The goal is to:
 - 1 learn w_0 and w_1
 - 2 predict the value of y based on the value of x

The best fitting line (regression line)

- Linear regression can be understood as finding the best-fitting straight line through the sample points, as shown in Figure 1 (see next page)
- The best-fitting line is also called the regression line
- The vertical lines from the regression line to the sample points are the errors of our prediction (offsets, or residuals)

Figure 1



Multiple linear regression

- We can generalize simple linear regression with only one explanatory variable to a model with multiple explanatory variables:

$$y = w_0x_0 + w_1x_1 + \cdots + w_mx_m = \sum_{i=0}^m w_ix_i = \mathbf{w}^T \mathbf{x}$$

- Such model is called multiple linear regression

Visualizing the important characteristics of a dataset

- Exploratory Data Analysis (EDA) is an important and recommended first step prior to the training of a machine learning model
- The graphical EDA toolbox may help
 - visually detect the presence of outliers
 - the distribution of the data
 - the relationship between features

Two kinds of useful graphical summaries

- Scatterplot matrix: allows us to visualize the pair-wise correlations between different features
- Correlation matrix:
 - a square matrix that contains the Pearson product-moment correlation coefficients (or Pearson's r), which measures the linear dependence between pairs of features
 - can be calculated as the covariance between two features x and y , divided by the product of their standard deviations


$$r = \frac{\sum_{i=1}^n [(x^{(i)} - \mu_x)(y^{(i)} - \mu_y)]}{\sqrt{\sum_{i=1}^n (x^{(i)} - \mu_x)^2} \sqrt{\sum_{i=1}^n (y^{(i)} - \mu_y)^2}} = \frac{\sigma_{xy}}{\sigma_x \sigma_y}$$

- identical to a covariance matrix computed from standardized data
- Combine the two summaries for choosing explanatory variables

Evaluating the performance using residual plot

- When our model uses multiple explanatory variables, we cannot visualize the model in a two-dimensional plot
- Instead we can plot the residuals (the difference or vertical distances between the actual and predicted values) versus the predicted values to diagnose our regression model
- One residual plot is shown in Figure 2 (see next page)

Figure 2



Evaluating the performance using Mean Squared Error

- Another useful quantitative measure of a model's performance is the so-called Mean Squared Error (MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

- The MSE is useful for comparing different regression models or for tuning their parameters via a grid search and cross-validation
- See details in `ch10.ipynb`

Estimating the parameters with gradient descent

- The parameters of the regression can be approximated by minimizing the cost function
- The cost function here is the Ordinary Least Squares (OLS) function

$$J(w) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

- We can minimize the cost function to learn the weights via optimization algorithms such as Gradient Descent (GD)
- Using GD, the rule for updating the weights can be written as

$$\mathbf{w} = \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla J(\mathbf{w}).$$

Here:

- η is the learning rate
- $\nabla J(\mathbf{w})$ the gradient of $J(\mathbf{w})$
- **Q:** Why the negative sign?

Estimating the parameters with gradient descent

- The parameters of the regression can be approximated by minimizing the cost function
- The cost function here is the Ordinary Least Squares (OLS) function

$$J(w) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

- We can minimize the cost function to learn the weights via optimization algorithms such as Gradient Descent (GD)
- Using GD, the rule for updating the weights can be written as

$$\mathbf{w} = \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla J(\mathbf{w}).$$


Here:

- η is the learning rate
- $\nabla J(\mathbf{w})$ the gradient of $J(\mathbf{w})$
- **Q:** Why the negative sign?
- **A:** Two kinds of explanations

The visual explanation

- The negative sign updates the weights by taking a step in the opposite direction of the gradient
- Thus we can minimize, rather than maximize, the cost function
- This is also the reason why it is called gradient descent, rather than gradient ascent
- Figure 3 (see next page) shows the idea

Figure 3



The mathematical explanation

- The updating rule for the weight of x_j , w_j , is

$$w_j = w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = -\eta \frac{\partial J}{\partial w_j}.$$

Here, $\frac{\partial J}{\partial w_j}$ is the partial derivative of J with respect to w_j :

$$\frac{\partial J}{\partial w_j} = - \sum_i (y^{(i)} - \hat{y}^{(i)}) x_j^i.$$

- Thus, the updating rule for w_j can be written as

$$w_j = w_j + \eta \sum_i (y^{(i)} - \hat{y}^{(i)}) x_j^i.$$

The mathematical explanation

- **Q:** Why does the updating rule work:

$$w_j = w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = \eta \sum_i (y^{(i)} - \hat{y}^{(i)}) x_j^i$$

- **A:** Because it pulls the predicted value (\hat{y}) closer to the actual one (y)
- Assume $y^{(i)} = 1$, $\hat{y}^{(i)} = -1$, and $x_j^{(i)} > 0$, then:
 - ① $\Delta w_j > 0$
 - ② $w_j \uparrow$
 - ③ $\hat{y} = \mathbf{w}^T \mathbf{x} \uparrow$

Fitting a robust regression model using RANSAC

- Linear regression models can be heavily impacted by the outliers
- In certain situations, a very small subset of our data can have a big effect on the estimated model coefficients (the parameters)
- RANdom SAmple Consensus (RANSAC) algorithm is an alternative to throwing out outliers, by fitting a regression model to a subset of the data, the so-called inliers

The RANSAC algorithm

- Basic steps in RANSAC
 - ① Select a random number of samples to be inliers and fit the model
 - ② Test all other data points against the fitted model and add those points that fall within a user-given tolerance to the inliers
 - ③ Refit the model using all inliers
 - ④ Estimate the error of the fitted model versus the inliers
 - ⑤ Terminate the algorithm if the performance meets a certain user-defined threshold or if a fixed number of iterations has been reached; go back to step 1 otherwise
- See details in `ch10.ipynb`