

Machine Learning I (DATS 6202)


Logistic Regression

Yuxiao Huang

Data Science, Columbian College of Arts & Sciences
George Washington University
yuxiaohuang@gwu.edu

September 18, 2018

Reference



Picture courtesy of the website of the book code repository and info resource

Reference

- This set of slides is an excerpt of the book by Raschka and Mirjalili, with some trivial changes by the creator of the slides
- Please find the reference to and website of the book below:
 - *Raschka S. and Mirjalili V. (2017). Python Machine Learning. 2nd Edition.*
 - <https://sebastianraschka.com/books.html>
- Please find the website of the book code repository and info resource below:
 - <https://github.com/rasbt/python-machine-learning-book-2nd-edition>

Overview

- 1 Motivation
- 2 Intuition
- 3 The logistic regression model
- 4 Training a logistic regression model
- 5 Tackling overfitting via regularization

Example: applying for a credit card

Credit Score	Approve
750	yes
700	no
...	...
650	yes
760	yes

Example: applying for a credit card

Credit Score	Approve
750	yes
700	no
...	...
650	yes
760	yes
680	?

Example: applying for a credit card

Credit Score	Approve
750	yes
700	no
...	...
650	yes
760	yes
680	$P(\text{yes} \text{680})$

Problem statement

- Given:
 - Credit score: x
 - Approve: y
- Predict:
 - Probability of y given x : $P(y|x)$

Linear regression?

- Linearity assumption:

$$P(y|x) = w_0 + w_1x$$

Linear regression?

- Linearity assumption:

$$P(y|x) = w_0 + w_1x$$

- Any problems?

Linear regression?

- Linearity assumption:

$$P(y|x) = w_0 + w_1x$$

- Any problems?

$$\underbrace{P(y|x)}_{[0,1]} = \underbrace{w_0 + w_1x}_{[-\infty, +\infty]}$$

Solution

- Linearity assumption:

$$\underbrace{P(y|x)}_{[0,1]} = \underbrace{w_0 + w_1 x}_{[-\infty, +\infty]}$$

Solution

- Linearity assumption:

$$\underbrace{P(y|x)}_{[0,1]} = \underbrace{w_0 + w_1 x}_{[-\infty, +\infty]}$$

- Find a function f such that

$$\underbrace{f(P(y|x))}_{[-\infty, +\infty]} = \underbrace{w_0 + w_1 x}_{[-\infty, +\infty]}$$

The *odds* function

- Find a function f such that


$$\underbrace{f(P)}_{[-\infty, +\infty]} = \underbrace{w_0 + w_1 x}_{[-\infty, +\infty]}$$

- The *odds* function:

$$\text{odds}(P) = \underbrace{\frac{P}{1-P}}_{[0, +\infty]}$$

- What else can we do?

Figure 1



The *log odds* function

- Find a function f such that

$$\underbrace{f(P)}_{[-\infty, +\infty]} = \underbrace{w_0 + w_1 x}_{[-\infty, +\infty]}$$

- The *odds* function:

$$\text{odds}(P) = \underbrace{\frac{P}{1-P}}_{[0, +\infty]}$$

- The *log odds* function

$$\text{log odds}(P) = \underbrace{\log\left(\frac{P}{1-P}\right)}_{[-\infty, +\infty]}$$

The *logit* function

- Find a function f such that

$$\underbrace{f(P)}_{[-\infty, +\infty]} = \underbrace{w_0 + w_1 x}_{[-\infty, +\infty]}$$

- The *odds* function:

$$\text{odds}(P) = \frac{P}{\underbrace{1 - P}_{[0, +\infty]}}$$

- The $\underbrace{\log \text{ odds}}_{\text{logit}}$ function

$$\underbrace{\log \text{ odds}(P)}_{\text{logit}} = \underbrace{\log \left(\frac{P}{1 - P} \right)}_{[-\infty, +\infty]}$$

The *logit* function

- The *logit* function takes input values in the range 0 to 1 and transforms them to values over the entire real number range

$$\underbrace{\text{logit}(P(y|x))}_{[0,1]} = \underbrace{\log\left(\frac{P(y|x)}{1 - P(y|x)}\right)}_{[-\infty, +\infty]}$$

- Here $P(y|x)$ is the probability that a certain sample belongs to a particular class

Logistic regression


- Logistic regression expresses a linear relationship between the *logit* and feature values

$$\underbrace{\text{logit}(P(y|x))}_{[-\infty, +\infty]} = \underbrace{w_0x_0 + w_1x_1 + \cdots + x_mw_m}_{[-\infty, +\infty]} = \sum_{i=0}^m w_ix_i = \mathbf{w}^T \mathbf{x}$$

- The goal is predicting $P(y|x)$ using the inverse of the *logit* function

$$P(y|x) = \text{logit}^{-1}(\mathbf{w}^T \mathbf{x})$$

Figure 2



The *logistic* function

- The inverse of the *logit* function, logit^{-1} , is also called the *logistic* function, sometimes simply abbreviated as the *sigmoid* function (due to its characteristic S-shape), shown in Figure 2 (see previous page)

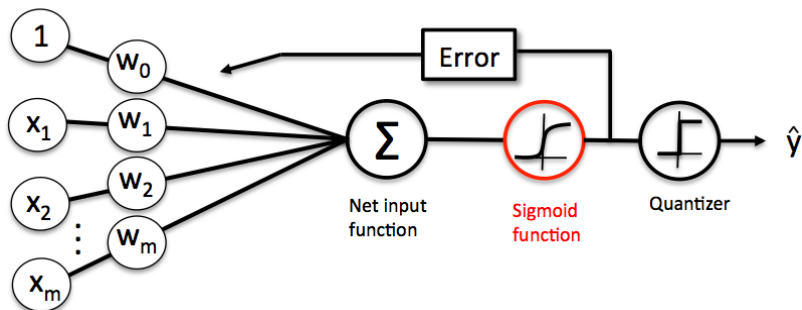
$$\phi(z) = \frac{1}{1 + e^{-z}}$$

- Here, z is the net input, that is, the linear combination of weights and sample features and can be calculated as

$$z = w_0x_0 + w_1x_1 + \cdots + x_mw_m = \sum_{i=0}^m w_ix_i = \mathbf{w}^T \mathbf{x}$$

- The logistic regression model is shown in Figure 3 (see next page)
- See details in `ch3.ipynb`

Figure 3



The quantizer

- The predicted probability (the outcome of the sigmoid function) can simply be converted into a binary outcome via a quantizer (unit step function)

$$\hat{y} = \begin{cases} 1 & \text{if } \phi(z) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

- This is equivalent to the following

$$\hat{y} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

- In fact, there are many applications where we are interested in not only the predicted class labels, but also the probability

Learning the weights of the logistic cost function

- The joint likelihood $L(D|\mathbf{w})$, assuming that the individual likelihood of each sample are independent, is

$$L(D|\mathbf{w}) = \prod_{i=1}^n P(y^{(i)}|x^{(i)}; \mathbf{w}) = \prod_{i=1}^n \left(\phi(z^{(i)}) \right)^{y^{(i)}} \left(1 - \phi(z^{(i)}) \right)^{1-y^{(i)}}$$

- The parameters, \mathbf{w} , can then be estimated by maximizing the joint likelihood, $L(D|\mathbf{w})$:

$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} L(D|\mathbf{w})$$

- This approach is called Maximum Likelihood Estimation (MLE)
- MLE can be solved using optimization algorithms such as gradient ascent (or gradient descent when rewriting the likelihood as a cost function and minimizing it)

The log trick

- In practice, it is easier to maximize the (natural) log of this equation, which is called the log-likelihood function:

$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} (\log L(D|\mathbf{w}))$$

- Here the log-likelihood function is

$$\log (L(D|\mathbf{w})) = \sum_{i=1}^n \left[y^{(i)} \log \left(\phi(z^{(i)}) \right) + \left(1 - y^{(i)} \right) \log \left(1 - \phi(z^{(i)}) \right) \right]$$

- The log trick can
 - reduce the potential for numerical underflow (when the likelihoods are very small)
 - simplify the derivation (by converting the product of factors into a summation of factors)

Estimating the parameters with gradient descent

- The parameters of the model can be approximated by minimizing the cost function (the additive inverse of the log-likelihood)

$$\begin{aligned} J(\mathbf{w}) &= -\log(L(D|\mathbf{w})) \\ &= -\sum_{i=1}^n \left[y^{(i)} \log \left(\phi(z^{(i)}) \right) + \left(1 - y^{(i)} \right) \log \left(1 - \phi(z^{(i)}) \right) \right] \end{aligned}$$

- As in linear regression, here we can minimize the cost function to learn the weights via Gradient Descent (GD)
- Using GD, the rule for updating the weights can be written as

$$\mathbf{w} = \mathbf{w} + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = -\eta \nabla J(\mathbf{w}).$$

Here:

- η is the learning rate
- $\nabla J(\mathbf{w})$ the gradient of $J(\mathbf{w})$

The updating rule

- The rule for updating w_j can be written as

$$w_j = w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = \eta \sum_i \left(y^{(i)} - \phi(z^{(i)}) \right) x_j^i.$$

- **Q:** Why does the updating rule work?

The updating rule

- The rule for updating w_j can be written as

$$w_j = w_j + \Delta w_j \quad \text{where} \quad \Delta w_j = \eta \sum_i \left(y^{(i)} - \phi(z^{(i)}) \right) x_j^i.$$

- Q:** Why does the updating rule work?
- A:** Because it pulls the predicted probability ($\phi(z)$) closer to the actual one (y)
- Assume $y^{(i)} = 1$, $\phi(z^{(i)}) = 0.7$, and $x_j^{(i)} > 0$, then:
 - $\Delta w_j > 0$
 - $w_j \uparrow$
 - $\phi(z) \uparrow$


Training a logistic regression model with scikit-learn

- See details in `ch3.ipynb`

Overfitting and underfitting

- **Overfitting** is a common problem in machine learning, where a model performs well on training data but does not generalize well to unseen data (test data)
- If a model suffers from overfitting, we also say that the model has a high variance, which can be caused by having too many parameters that lead to a model that is too complex given the underlying data
- Similarly, our model can also suffer from **underfitting** (high bias), which means that our model is not complex enough to capture the pattern in the training data well and therefore also suffers from low performance on unseen data
- Figure 4 illustrates overfitting and underfitting (see next page)

Figure 4



Regularization

- One way of finding a good bias-variance tradeoff is to tune the complexity of the model via regularization
- Regularization is a very useful method to handle collinearity (high correlation among features), filter out noise from data, and eventually prevent overfitting
- The idea behind regularization is to introduce additional information (bias) to penalize extreme parameter (weight) values

Regularization

- The most common form of regularization is so-called L2 regularization (sometimes also called L2 shrinkage or weight decay), which can be written as follows:

$$\frac{\lambda}{2} \|\mathbf{w}\|^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2 \quad (1)$$

where λ is the so-called regularization parameter

- The cost function for logistic regression can be regularized by adding a simple regularization term, which will shrink the weights during model training:

$$J(\mathbf{w}) = - \sum_{i=1}^n \left[y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right] + \frac{\lambda}{2} \|\mathbf{w}\|^2 \quad (2)$$

Regularization

- Via the regularization parameter λ , we can then control how well we fit the training data while keeping the weights small
 - by increasing the value of λ , we increase the regularization strength
- The parameter `C` that is implemented for the `LogisticRegression` class in `scikit-learn` is the inverse of the regularization parameter λ
 - by decreasing the value of `C`, we increase the regularization strength
- See details in `ch3.ipynb`