

Machine Learning I (DATS 6202)

Shallow Neural Network

Yuxiao Huang

Data Science, Columbian College of Arts & Sciences
George Washington University

Spring 2020

Reference

- This set of slides was largely built on the following 8 wonderful books and a wide range of fabulous papers:
 - HML Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow (2nd Edition)
 - PML Python Machine Learning (3rd Edition)
 - ESL The Elements of Statistical Learning (2nd Edition)
 - PRML Pattern Recognition and Machine Learning
 - LFD Learning From Data
 - NND Neural Network Design (2nd Edition)
 - NNDL Neural Network and Deep Learning
 - RL Reinforcement Learning: An Introduction
- For most materials covered in the slides, we will specify their corresponding books and papers for further reference.

Code Example & Case Study

- See code example of topics discussed in this section in github repository: [/p1_c2_s7_shallow_neural_network/code_example](#)
- See case study of Kaggle Competition related to this section in github repository: [/p1_c2_s7_shallow_neural_network/case_study](#)

Overview

- 1 Learning Objectives
- 2 Single-Layer Perceptron
- 3 Perceptron Learning Rule
- 4 Multi-Layer Perceptron
- 5 Backpropagation
- 6 References

Learning Objectives

- It is **expected** to understand
 - the idea of single-layer perceptron and perceptron learning rule
 - the idea of multi-layer perceptron and backpropagation
 - the pros and cons of batch, stochastic and mini-batch gradient descent
 - how to use sklearn single-layer perceptron and multi-layer perceptron
- It is **recommended** to understand
 - the math behind perception learning rule and backpropagation
 - how to implement single-layer perceptron (using perception learning rule) and multi-layer perceptron (using backpropagation and mini-batch gradient descent)

Biological Neuron

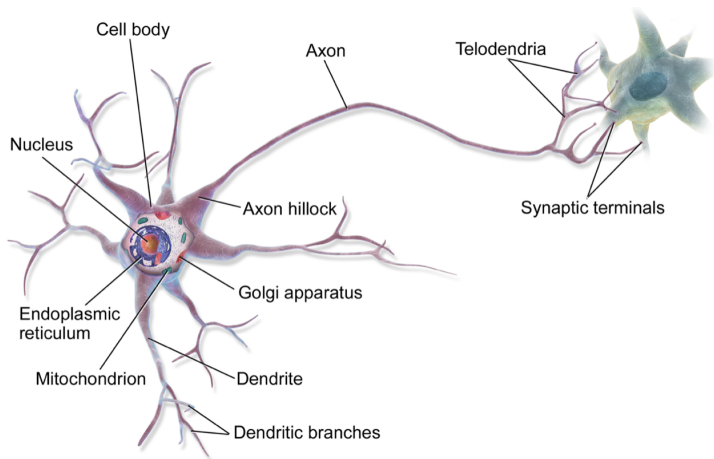


Figure 1: Biological neuron. Picture courtesy of *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow (2nd Edition)*.

Single-Layer Perceptron (SLP)

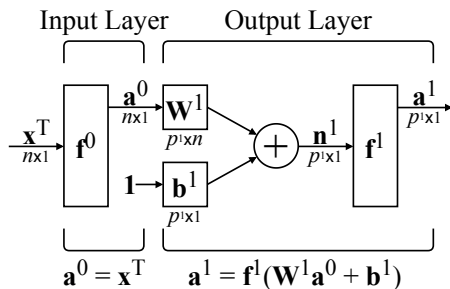


Figure 2: A single-layer perceptron.

Here:

- \mathbf{x}^T is the input of the input layer, with n features and 1 sample
- \mathbf{a}^1 is the output of the output layer, with p^1 perceptrons

Some Popular Activation Functions

- HardLimit:

$$HardLimit(n) = \begin{cases} 1, & \text{if } n \geq 0, \\ 0, & \text{otherwise.} \end{cases} \quad (1)$$

- Rectified Linear Unit (ReLU):

$$ReLU(n) = \begin{cases} n, & \text{if } n \geq 0, \\ 0, & \text{otherwise.} \end{cases} \quad (2)$$

- Linear:

$$Linear(n) = n. \quad (3)$$

- Sigmoid (σ):

$$\sigma(n) = \frac{1}{1 + e^{-n}}. \quad (4)$$

- Hyperbolic Tangent ($tanh$):

$$tanh(n) = 2\sigma(2n) - 1. \quad (5)$$

The Logic AND Data




Figure 3: The logic AND data.

A Corresponding SLP

- **Q:** Can you design a slp that perfectly separates the logic AND data?

A Corresponding SLP

- **Q:** Can you design a slp that perfectly separates the logic AND data?

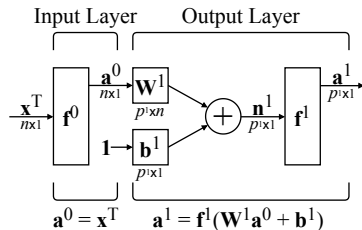


Figure 4: A single-layer perceptron.

- **A:** A slp in fig. 4 with the following parameter settings:

- $\mathbf{w}^1 = \begin{bmatrix} 1 & 1 \end{bmatrix}$
- $b^1 = -1.5$
- $f^1 = \text{HardLimit}$

The Logic OR Data

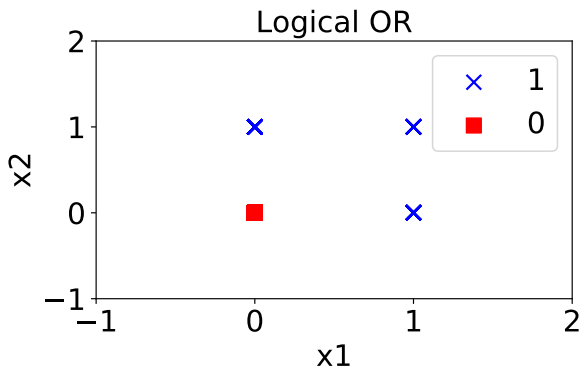


Figure 5: The logic OR data.

A Corresponding SLP

- **Q:** Can you design a slp that perfectly separates the logic OR data?

A Corresponding SLP

- **Q:** Can you design a slp that perfectly separates the logic OR data?

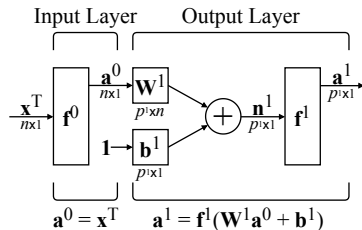


Figure 6: A single-layer perceptron.

- **A:** A slp in fig. 6 with the following parameter settings:
 - $\mathbf{w}^1 = \begin{bmatrix} 1 & 1 \end{bmatrix}$
 - $b^1 = -0.5$
 - $f^1 = \text{HardLimit}$

The Motivation

- As shown in the previous two problems (logic AND and logic OR), designing the parameters of a slp manually may not be easy.
- Would not it be great if we can learn the parameters of slp automatically, as what we did in linear and logistic regression?
- It turns out that we can use the *Perceptron Learning Rule* to do so.

Perceptron Learning Rule: The Idea + Math

- For each sample, the parameters of a slp are updated using the perceptron learning rule:

$$\mathbf{W}^1 = \mathbf{W}^1 + \eta(\mathbf{y} - \hat{\mathbf{y}})\mathbf{x}, \quad (6)$$

$$\mathbf{b}^1 = \mathbf{b}^1 + \eta(\mathbf{y} - \hat{\mathbf{y}}). \quad (7)$$

Here:

- \mathbf{W}^1 is the connecting weight between the output layer and input layer
- \mathbf{b}^1 is the bias of the output layer
- \mathbf{y} is the target of the current sample
- $\hat{\mathbf{y}}$ is the output of the output layer
- \mathbf{x} is the current sample

Further Reading

- See NND Chapter 4 for a very nice explanation of why the perceptron learning rule works, from a linear algebra perspective.
- See NND Chapter 4 for the proof of convergence for the perceptron learning rule (i.e., the proof that slp can always perfectly separate linearly-separable data).

The Logic XOR Data

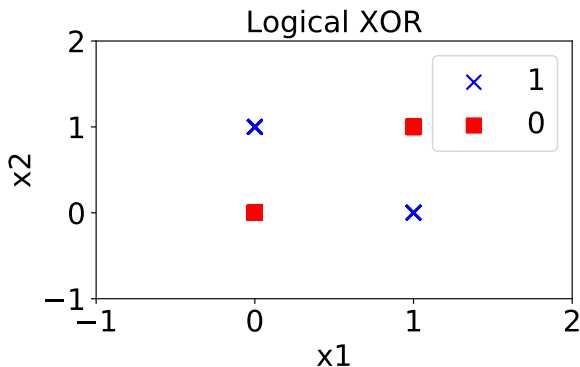


Figure 7: The logic XOR data.

A Corresponding SLP

- **Q:** Can you design a slp that perfectly separates the logic XOR data?

A Corresponding SLP

- **Q:** Can you design a slp that perfectly separates the logic XOR data?
- **A:** Unfortunately, no, since the data is not linearly-separable.

Biological Neural Network

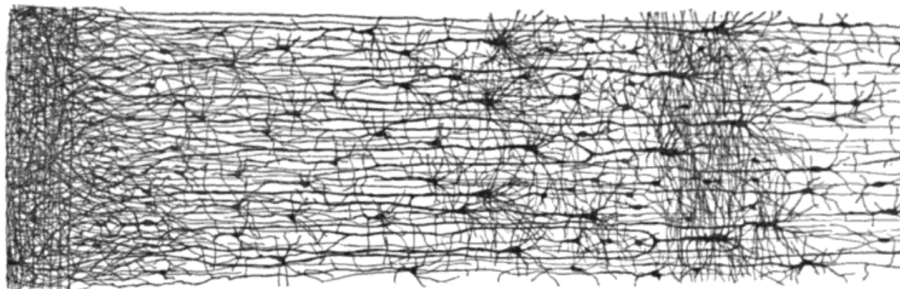


Figure 8: Biological neural network. Picture courtesy of *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow (2nd Edition)*.

Multi-Layer Perceptron (MLP)

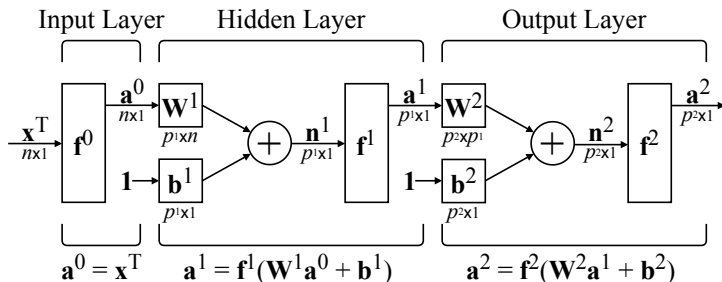


Figure 9: A fully connected three-layer mlp.

Here:

- \mathbf{x}^T is the input of the input layer, with n features and 1 sample
- \mathbf{a}^1 is the output of the hidden layer, with p^1 perceptrons
- \mathbf{a}^2 is the output of the output layer, with p^2 perceptrons

A MLP for the Logic XOR Data

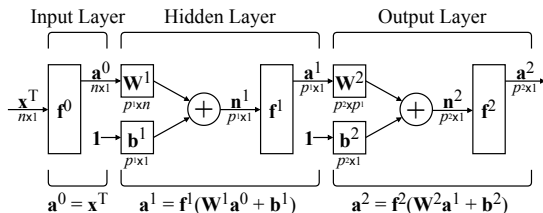


Figure 10: A fully connected three-layer mlp for the logic XOR data.

- A mlp in fig. 10 with the following parameter settings:

- $\mathbf{w}^1 = \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix}$, $\mathbf{b}^1 = \begin{bmatrix} -0.5 \\ -0.5 \end{bmatrix}$
- $\mathbf{w}^2 = \begin{bmatrix} 1 & 1 \end{bmatrix}$, $\mathbf{b}^2 = -0.5$
- $f^1 = f^2 = \text{HardLimit}$

MLP with Multiple Hidden Layers

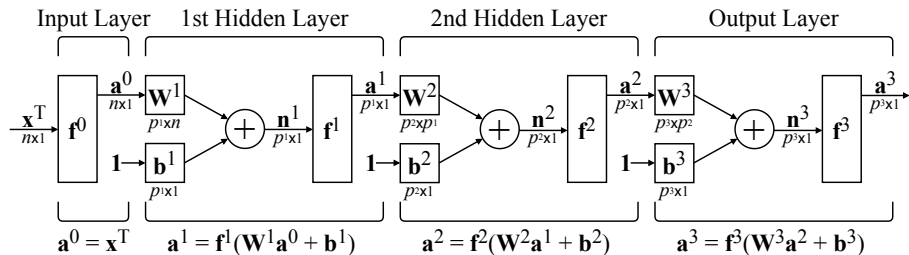


Figure 11: A fully connected four-layer mlp.

Here:

- \mathbf{x}^T is the input of the input layer, with n features and 1 sample
- \mathbf{a}^1 is the output of the first hidden layer, with p^1 perceptrons
- \mathbf{a}^2 is the output of the second hidden layer, with p^2 perceptrons
- \mathbf{a}^3 is the output of the output layer, with p^3 perceptrons

The Motivation

- Since manually designing the parameters of a slp is not easy, doing so for a mlp is much more difficult, if not impossible.
- Unfortunately, we cannot simply use perceptron learning rule to automatically train a mlp (we will see the reason later).
- Instead, we can use the *Backpropagation* to do so.

The Idea

- Backpropagation is an iterative process.
- There are three steps (in order) in each iteration:
 - ① forward pass: sends information from input layer up to higher layers
 - ② backward pass: sends information from output layer down to lower layers
 - ③ gradient descent: updates the parameters using information obtained in the forward and backward pass

Forward Pass: The Idea

- Forward pass works as follows:
 - ① passes the input to the input layer
 - ② from the first hidden layer up to the output layer, calculates the output of the current layer (e.g., layer k) from the output of the previous layer (e.g., layer $k - 1$)
- In forward pass, the net input and activation of each layer are preserved and reused in backward pass and gradient descent.

Backward Pass: The Idea

- Backward pass works as follows:
 - ① calculates the objective function (e.g., Mean Squared Error) based on the target and the output of the output layer
 - ② calculates the *Sensitivity* (more on this later) of the output layer
 - ③ from the last hidden layer down to the first hidden layer, calculates the sensitivity of the current layer (e.g., layer k) from the sensitivity of the next layer (e.g., layer $k + 1$)
- In backward pass, the sensitivity of each layer is preserved and reused in gradient descent

Gradient Descent: The Idea

- Gradient descent works as follows:
 - ① uses the output (obtained in forward pass) and sensitivity (obtained in backward pass) of each layer to update the parameters of the layer

The Math

- Let us take a look at the math underlying the three steps in backpropagation (forward pass, backward pass and gradient descent)
- The order in which the steps work:
 - ① forward pass
 - ② backward pass
 - ③ gradient descent
- The order in which we will discuss the steps (which better explains the dependence between the steps):
 - ① gradient descent
 - ② forward pass
 - ③ backward pass

Batch / Stochastic / Mini-Batch Gradient Descent

- Based on the amount of data that we use in each iteration of updating the parameters, we can divide gradient descent into three categories:
 - *batch gradient descent*: where we use all the training data
 - *stochastic gradient descent*: where we use a sample in the training data
 - *mini-batch gradient descent*: where we use a subset (with more than one sample) of the training data
- We can think of batch gradient descent / stochastic gradient descent as two extremes (using as many / as few data as possible).
- In that sense, mini-batch lies between the two extremes.

Batch Gradient Descent

- pros:
 - utilizes parallel computing to the most extent (making it fast)
 - does not require shuffling the data (making it even faster)
- cons:
 - requires loading all the training data into the memory (not suitable for large datasets)
 - requires processing all the training data to complete one update of the parameters (not suitable for large datasets)
- Batch gradient descent is suitable for small datasets.

Stochastic Gradient Descent

- pros:
 - allows loading only some training data into the memory (suitable for large datasets)
 - allows processing only one sample in the training data to complete one update of the parameters (suitable for large datasets)
- cons:
 - utilizes parallel computing to the least extent (making it slow)
 - requires shuffling the data in the beginning of each epoch (making it even slower)
- Stochastic gradient descent is suitable for large datasets.

Mini-Batch Gradient Descent

- pros:
 - allows loading only some training data into the memory (suitable for large datasets)
 - allows processing only some training data to complete one update of the parameters (suitable for large datasets)
- cons:
 - utilizes parallel computing to some extent (making it faster than stochastic gradient descent but slower than batch gradient descent)
 - requires shuffling the data in the beginning of each epoch (making it even slower than batch gradient descent)
- Mini-batch gradient descent is more suitable for large datasets (compared to stochastic gradient descent).
- **Good Practice:** It is recommended to use small batches (from 2 to 32) in mini-batch gradient descent [1].

The Objective Function

- While mini-batch gradient descent usually works better for large datasets, here we use stochastic gradient descent for mathematical convenience.
- Since stochastic gradient descent processes one sample at a time, the objective function (Mean Squared Error), J , can be written as

$$J = (\mathbf{y} - \hat{\mathbf{y}})^2. \quad (8)$$

Here $\hat{\mathbf{y}}$ is the output of the output layer:

$$\hat{\mathbf{y}} = \mathbf{a}^l = \mathbf{f}^l(\mathbf{n}^l) = \mathbf{f}^l(\mathbf{W}^l \mathbf{a}^{l-1} + \mathbf{b}^l), \quad (9)$$

where l is the index of the last layer and \mathbf{a}^l the output of the output layer.

Gradient Descent: The Math

- Using gradient descent, the weight with respect to layer k , \mathbf{W}^k , can be updated as

$$\mathbf{W}^k = \mathbf{W}^k - \eta \frac{\partial J}{\partial \mathbf{W}^k}. \quad (10)$$

- We cannot directly calculate the gradient in eq. (10) since, as shown in eqs. (8) and (9), while J is an explicit function of weights on the output layer, \mathbf{W}^l , it is not an explicit function of weights on the hidden layers, \mathbf{W}^k (where $k < l$).
- This is why we cannot use the perceptron learning rule for mlp, as we did for slp (where there is no hidden layer).
- This is also why we use the chain rule to rewrite eq. (10):

$$\mathbf{W}^k = \mathbf{W}^k - \eta \frac{\partial J}{\partial \mathbf{n}^k} \frac{\partial \mathbf{n}^k}{\partial \mathbf{W}^k}, \quad (11)$$

where \mathbf{n}^k is the net input of layer k (and an explicit function of \mathbf{W}^k):

$$\mathbf{n}^k = \mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k. \quad (12)$$

Gradient Descent: The Math

- We define the first derivation in eq. (11), $\frac{\partial J}{\partial \mathbf{n}^k}$, as the *Sensitivity* of J to the changes of \mathbf{n}^k , denoted by \mathbf{s}^k :

$$\mathbf{s}^k = \frac{\partial J}{\partial \mathbf{n}^k}. \quad (13)$$

- Based on eq. (12), we rewrite the second derivation in eq. (11), $\frac{\partial \mathbf{n}^k}{\partial \mathbf{W}^k}$, as

$$\frac{\partial \mathbf{n}^k}{\partial \mathbf{W}^k} = \frac{\partial (\mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k)}{\partial \mathbf{W}^k} = \left(\mathbf{a}^{k-1} \right)^\top. \quad (14)$$

Gradient Descent: The Math

- By substituting eqs. (13) and (14) into eq. (11)

$$\mathbf{W}^k = \mathbf{W}^k - \eta \frac{\partial J}{\partial \mathbf{n}^k} \frac{\partial \mathbf{n}^k}{\partial \mathbf{W}^k},$$

we can update the weight as

$$\mathbf{W}^k = \mathbf{W}^k - \eta \mathbf{s}^k \left(\mathbf{a}^{k-1} \right)^\top. \quad (15)$$

- Similarly, we can update the bias as

$$\mathbf{b}^k = \mathbf{b}^k - \eta \mathbf{s}^k. \quad (16)$$

- We use forward pass to calculate \mathbf{a}^{k-1} in eq. (15).
- We use backward pass to calculate \mathbf{s}^k in eqs. (15) and (16).

Gradient Descent: The Idea + Math

- Gradient descent works as follows:
 - ① uses eq. (15), the output (obtained in forward pass) and sensitivity (obtained in backward pass) of each layer to update the parameters of the layer:

$$\begin{aligned}\mathbf{W}^k &= \mathbf{W}^k - \eta \mathbf{s}^k (\mathbf{a}^{k-1})^\top, \\ \mathbf{b}^k &= \mathbf{b}^k - \eta \mathbf{s}^k.\end{aligned}$$

Forward Pass: The Math

- The key in forward pass is a recursive relationship between the output of two consecutive layers.
- Since the output of layer k , \mathbf{a}^k , is

$$\mathbf{a}^k = \mathbf{f}^k(\mathbf{n}^k), \quad (17)$$

where \mathbf{f}^k is the activation function of layer k and \mathbf{n}^k the net input of layer k , given in eq. (12)

$$\mathbf{n}^k = \mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k,$$

we can rewrite eq. (17) as

$$\mathbf{a}^k = \mathbf{f}^k(\mathbf{n}^k) = \mathbf{f}^k(\mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k). \quad (18)$$

- Eq. (18) shows that the output of layer k , \mathbf{a}^k , relies on the output of layer $k - 1$, \mathbf{a}^{k-1} .
- This is why we use *forward* pass to calculate the output from lower layers (e.g., layer $k - 1$) up to higher layers (e.g., layer k).

Forward Pass: The Idea + Math

- Forward pass works as follows:
 - ① passes the input to the input layer:

$$\mathbf{a}^0 = \mathbf{x}^\top.$$

- ② from the first hidden layer (layer 1) up to the output layer (layer l), uses eq. (18) to calculate the output of the current layer (e.g., layer k) from the output of the previous layer (e.g., layer $k - 1$):

$$\mathbf{a}^k = \mathbf{f}^k(\mathbf{n}^k) = \mathbf{f}^k(\mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k), \quad \text{where } 1 \leq k \leq l.$$

Backward Pass: The Math

- Similar to forward pass, the key in backward pass is also a recursive relationship.
- Unlike forward pass where the relationship is between *output* of two consecutive layers, the relationship in backward pass is between *sensitivity* of two consecutive layers.
- Concretely, by applying the chain rule to eq. (13)

$$\mathbf{s}^k = \frac{\partial J}{\partial \mathbf{n}^k},$$

we have

$$\mathbf{s}^k = \frac{\partial J}{\partial \mathbf{n}^k} = \left(\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} \right)^\top \frac{\partial J}{\partial \mathbf{n}^{k+1}} = \left(\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} \right)^\top \mathbf{s}^{k+1}. \quad (19)$$

- Eq. (19) shows that the sensitivity of layer k , \mathbf{s}^k , relies on the sensitivity of layer $k + 1$, \mathbf{s}^{k+1} .
- This is why we use *backward* pass to calculate the sensitivity from higher layers (e.g., layer $k + 1$) down to lower layers (e.g., layer k).

The Sensitivity

- Based on eq. (13)

$$\mathbf{s}^k = \frac{\partial J}{\partial \mathbf{n}^k},$$

the sensitivity of the output layer (with index l), \mathbf{s}^l , is

$$\mathbf{s}^l = \frac{\partial J}{\partial \mathbf{n}^l}. \quad (20)$$

- By substituting eqs. (8) and (9)

$$J = (\mathbf{y} - \hat{\mathbf{y}})^2 \quad \text{and} \quad \hat{\mathbf{y}} = \mathbf{a}^l = \mathbf{f}^l(\mathbf{n}^l)$$

into eq. (20), we have

$$\mathbf{s}^l = \frac{\partial J}{\partial \mathbf{n}^l} = \frac{\partial (\mathbf{y} - \hat{\mathbf{y}})^2}{\partial \mathbf{n}^l} = -2(\mathbf{y} - \hat{\mathbf{y}}) \frac{\partial \mathbf{f}^l(\mathbf{n}^l)}{\partial \mathbf{n}^l} = -2(\mathbf{y} - \hat{\mathbf{y}}) \dot{\mathbf{f}}^l(\mathbf{n}^l). \quad (21)$$

The Jacobian Matrix

- The last piece in backward pass is the derivation in eq. (19)

$$\mathbf{s}^k = \left(\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} \right)^\top \mathbf{s}^{k+1}.$$

The transpose of the derivation is a Jacobian matrix:

$$\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} = \begin{bmatrix} \frac{\partial n_1^{k+1}}{\partial n_1^k} & \frac{\partial n_1^{k+1}}{\partial n_2^k} & \cdots & \frac{\partial n_1^{k+1}}{\partial n_{p^k}^k} \\ \frac{\partial n_2^{k+1}}{\partial n_1^k} & \frac{\partial n_2^{k+1}}{\partial n_2^k} & \cdots & \frac{\partial n_2^{k+1}}{\partial n_{p^k}^k} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial n_{p^{k+1}}^{k+1}}{\partial n_1^k} & \frac{\partial n_{p^{k+1}}^{k+1}}{\partial n_2^k} & \cdots & \frac{\partial n_{p^{k+1}}^{k+1}}{\partial n_{p^k}^k} \end{bmatrix}. \quad (22)$$

Here:

- entry (i, j) in the matrix is $\frac{\partial n_i^{k+1}}{\partial n_j^k}$
- n_i^{k+1} is the net input of the i^{th} perceptron on layer $k + 1$
- n_j^k is the net input of the j^{th} perceptron on layer k
- p^{k+1} and p^k are the number of perceptrons on layer $k + 1$ and k

The Jacobian Matrix

- In order to derive the equation of entry (i, j) in the Jacobian matrix, $\frac{\partial n_i^{k+1}}{\partial n_j^k}$, we first derive the equation of the net input of the i^{th} perceptron on layer $k + 1$, n_i^{k+1} .
- Since the i^{th} perceptron on layer $k + 1$ is fully connected with all the p^k perceptrons on layer k , we can write n_i^{k+1} as

$$n_i^{k+1} = \sum_{q=1}^{p^k} w_{iq}^{k+1} a_q^k + b_i^{k+1}. \quad (23)$$

Here:

- w_{iq}^{k+1} is the connecting weight between the i^{th} perceptron on layer $k + 1$ and the q^{th} perceptron on layer k
- a_q^k is the output of the q^{th} perceptron on layer k
- b_i^{k+1} is the bias of the i^{th} perceptron on layer $k + 1$

The Jacobian Matrix

- By substituting eq. (23) into $\frac{\partial n_i^{k+1}}{\partial n_j^k}$, we have

$$\frac{\partial n_i^{k+1}}{\partial n_j^k} = \frac{\partial \left(\sum_{q=1}^{p^k} w_{iq}^{k+1} a_q^k + b_i^{k+1} \right)}{\partial n_j^k}. \quad (24)$$

- Since b_i^{k+1} (the bias of the i^{th} perceptron on layer $k+1$) is a scalar, we can rewrite eq. (24) as

$$\frac{\partial n_i^{k+1}}{\partial n_j^k} = \frac{\partial \sum_{q=1}^{p^k} w_{iq}^{k+1} a_q^k}{\partial n_j^k}. \quad (25)$$

The Jacobian Matrix

- The output of the q^{th} perceptron on layer k , a_q^k , is

$$a_q^k = f_q^k(n_q^k), \quad (26)$$

where f_q^k and n_q^k are the activation function and net input of the q^{th} perceptron on layer k .

- Based on eq. (26), a_q^k is only related to n_j^k when $q = j$, we can rewrite eq. (25) as

$$\frac{\partial n_i^{k+1}}{\partial n_j^k} = \frac{\partial w_{ij}^{k+1} a_j^k}{\partial n_j^k} = w_{ij}^{k+1} \frac{\partial f_j^k(n_j^k)}{\partial n_j^k} = w_{ij}^{k+1} f_j^k(n_j^k). \quad (27)$$

The Jacobian Matrix

- By substituting eq. (27) into the Jacobian matrix in eq. (22)

$$\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} = \begin{bmatrix} \frac{\partial n_1^{k+1}}{\partial n_1^k} & \frac{\partial n_1^{k+1}}{\partial n_2^k} & \cdots & \frac{\partial n_1^{k+1}}{\partial n_{p^k}^k} \\ \frac{\partial n_2^{k+1}}{\partial n_1^k} & \frac{\partial n_2^{k+1}}{\partial n_2^k} & \cdots & \frac{\partial n_2^{k+1}}{\partial n_{p^k}^k} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial n_{p^{k+1}}^{k+1}}{\partial n_1^k} & \frac{\partial n_{p^{k+1}}^{k+1}}{\partial n_2^k} & \cdots & \frac{\partial n_{p^{k+1}}^{k+1}}{\partial n_{p^k}^k} \end{bmatrix},$$

we have

$$\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} = \mathbf{W}^{k+1} \dot{\mathbf{F}}^k(\mathbf{n}^k), \quad (28)$$

where $\dot{\mathbf{F}}^k(\mathbf{n}^k)$ is a diagonal matrix:

$$\dot{\mathbf{F}}^k(\mathbf{n}^k) = \begin{bmatrix} \dot{f}_1^k(n_1^k) & 0 & \cdots & 0 \\ 0 & \dot{f}_2^k(n_2^k) & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \dot{f}_{p^k}^k(n_{p^k}^k) \end{bmatrix}. \quad (29)$$

The Jacobian Matrix

- By substituting eq. (28) into eq. (19)

$$\mathbf{s}^k = \left(\frac{\partial \mathbf{n}^{k+1}}{\partial \mathbf{n}^k} \right)^\top \mathbf{s}^{k+1},$$

we have

$$\mathbf{s}^k = \dot{\mathbf{F}}^k(\mathbf{n}^k) \left(\mathbf{W}^{k+1} \right)^\top \mathbf{s}^{k+1}, \quad (30)$$

where $\dot{\mathbf{F}}^k(\mathbf{n}^k)$ is given in eq. (29)

$$\dot{\mathbf{F}}^k(\mathbf{n}^k) = \begin{bmatrix} \dot{f}_1^k(n_1^k) & 0 & \cdots & 0 \\ 0 & \dot{f}_2^k(n_2^k) & \cdots & 0 \\ \vdots & \vdots & & \vdots \\ 0 & 0 & \cdots & \dot{f}_{p^k}^k(n_{p^k}^k) \end{bmatrix}.$$

Backward Pass: The Idea + Math

- Backward pass works as follows:

- 1 calculates the objective function (Mean Squared Error) based on the target and the output of the output layer using eq. (8):

$$J = (\mathbf{y} - \hat{\mathbf{y}})^2.$$

- 2 calculates the sensitivity of the output layer (i.e., the last layer with index l) using eq. (21):

$$\mathbf{s}^l = \frac{\partial J}{\partial \mathbf{n}^l} = -2(\mathbf{y} - \hat{\mathbf{y}})\dot{\mathbf{f}}^l(\mathbf{n}^l).$$

- 3 from the last hidden layer down to the first hidden layer, uses eq. (30) to calculate the sensitivity of the current layer (e.g., layer k) from the sensitivity of the next layer (e.g., layer $k + 1$):

$$\mathbf{s}^k = \dot{\mathbf{F}}^k(\mathbf{n}^k) (\mathbf{W}^{k+1})^\top \mathbf{s}^{k+1}, \quad \text{where } l - 1 \geq k \geq 1.$$

Summary

① Forward pass

- ① passes the input to the input layer:

$$\mathbf{a}^0 = \mathbf{x}^\top.$$

- ② calculates the output from the first hidden layer up to the output layer:

$$\mathbf{a}^k = \mathbf{f}^k(\mathbf{n}^k) = \mathbf{f}^k(\mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k), \quad \text{where } 1 \leq k \leq l.$$

② Backward pass

- ① calculates the sensitivity of the output layer:

$$\mathbf{s}^l = \frac{\partial J}{\partial \mathbf{n}^l} = -2(\mathbf{y} - \hat{\mathbf{y}}) \dot{\mathbf{f}}^l(\mathbf{n}^l).$$

- ② calculates the sensitivity from the last hidden layer down to the first:

$$\mathbf{s}^k = \dot{\mathbf{F}}^k(\mathbf{n}^k) (\mathbf{W}^{k+1})^\top \mathbf{s}^{k+1}, \quad \text{where } l-1 \geq k \geq 1.$$

③ Gradient descent

- ① updates the parameters of each layer:

$$\begin{aligned} \mathbf{W}^k &= \mathbf{W}^k - \eta \mathbf{s}^k (\mathbf{a}^{k-1})^\top, \\ \mathbf{b}^k &= \mathbf{b}^k - \eta \mathbf{s}^k. \end{aligned}$$

Example

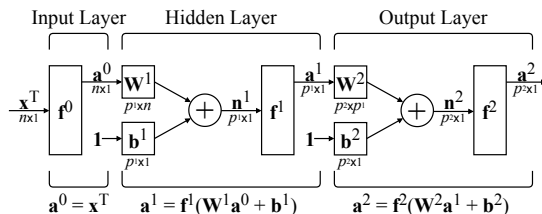


Figure 12: A fully connected three-layer mlp.

- Here is the first iteration of backpropagation on mlp in fig. 12, where:

- $\mathbf{x}^T = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$, $y = 0$, $n = 2$
- $\mathbf{W}^1 = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$, $\mathbf{b}^1 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$, $p^1 = 2$
- $\mathbf{W}^2 = \begin{bmatrix} 1 & 1 \end{bmatrix}$, $\mathbf{b}^2 = 0$, $p^2 = 1$
- $\mathbf{f}^1 = \text{ReLU}$, $\mathbf{f}^2 = \text{Linear}$

Forward Pass

- Passes the input to the input layer:

$$\mathbf{a}^0 = \mathbf{x}^T = \begin{bmatrix} 1 \\ 1 \end{bmatrix}. \quad (31)$$

- Calculates the output of the hidden layer:

$$\mathbf{n}^1 = \mathbf{W}^1 \mathbf{a}^0 + \mathbf{b}^1 = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 2 \\ 2 \end{bmatrix}, \quad (32)$$

$$\mathbf{a}^1 = \mathbf{f}^1(\mathbf{n}^1) = \mathbf{ReLU} \left(\begin{bmatrix} 2 \\ 2 \end{bmatrix} \right) = \begin{bmatrix} 2 \\ 2 \end{bmatrix}. \quad (33)$$

- Calculates the output of the output layer:

$$n^2 = \mathbf{w}^2 \mathbf{a}^1 + b^2 = \begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 2 \end{bmatrix} + 0 = 4, \quad (34)$$

$$\hat{y} = a^2 = f^2(n^2) = \text{Linear}(4) = 4. \quad (35)$$

Backward Pass

- Calculate the sensitivity of the output layer:

$$\dot{f}^2(n^2) = \text{Linear}(4) = 1, \quad (36)$$

$$s^2 = -2(y - \hat{y})\dot{f}^2(n^2) = -2 \times (0 - 4) \times 1 = 8. \quad (37)$$

- Calculate the sensitivity of the hidden layer:

$$\mathbf{\dot{F}}^1(\mathbf{n}^1) = \begin{bmatrix} \dot{f}_1^1(n_1^1) & 0 \\ 0 & \dot{f}_2^1(n_2^1) \end{bmatrix} = \begin{bmatrix} \text{ReLU}(2) & 0 \\ 0 & \text{ReLU}(2) \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad (38)$$

$$\mathbf{s}^1 = \mathbf{\dot{F}}^1(\mathbf{n}^1) (\mathbf{w}^2)^\top s^2 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} \times 8 = \begin{bmatrix} 8 \\ 8 \end{bmatrix}. \quad (39)$$

Gradient Descent

- Update the parameters of the hidden layer (with $\eta = 0.1$):

$$\mathbf{W}^1 = \mathbf{W}^1 - \eta \mathbf{s}^1 (\mathbf{a}^0)^\top = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} - 0.1 \times \begin{bmatrix} 8 \\ 8 \end{bmatrix} \begin{bmatrix} 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.2 & 0.2 \\ 0.2 & 0.2 \end{bmatrix}, \quad (40)$$

$$\mathbf{b}^1 = \mathbf{b}^1 - \eta \mathbf{s}^1 = \begin{bmatrix} 0 \\ 0 \end{bmatrix} - 0.1 \times \begin{bmatrix} 8 \\ 8 \end{bmatrix} = \begin{bmatrix} -0.8 \\ -0.8 \end{bmatrix}. \quad (41)$$

- Update the parameters of the output layer (with $\eta = 0.1$):

$$\mathbf{w}^2 = \mathbf{w}^2 - \eta s^2 (\mathbf{a}^1)^\top = \begin{bmatrix} 1 & 1 \end{bmatrix} - 0.1 \times 8 \times \begin{bmatrix} 2 & 2 \end{bmatrix} = \begin{bmatrix} -0.6 & -0.6 \end{bmatrix}, \quad (42)$$

$$b^2 = b^2 - \eta s^2 = 0 - 0.1 \times 8 = -0.8. \quad (43)$$

References



D. Masters and C. Lusch.

Revisiting Small Batch Training for Deep Neural Networks.

arXiv preprint [arXiv:1804.07612](https://arxiv.org/abs/1804.07612), 2018.